

=====

KRONUMOS AI (TOKENECTOMY LABS)
MICROSOFT FOR STARTUPS FOUNDERS HUB – SUBMISSION PACKAGE

=====

1. STARTUP & PRODUCT INFORMATION

- * Startup / Company Name:
Kronumos AI (Tokenectomy Labs)
- * Primary Website / Public Showcase:
<https://huggingface.co/Kronumos>
- * Open Source Repository:
<https://github.com/Tokenectomy-Labs/Kronomos>
- * Research Paper / Academic Provenance:
<https://doi.org/10.21203/rs.3.rs-11205335/v1> (Springer Nature Research Square)
- * Primary Country:
Indonesia
- * Current Funding Stage:
Bootstrapped / No outside funding yet

2. THREE CORE ESSAY QUESTIONS (COPY-PASTE READY)

[QUESTION A: What problem does your startup solve?]

Software engineering teams face prohibitive compute costs and high hallucination rates when using AI coding agents for automated software repair. Conventional agent loops consume 50,000 to 150,000 tokens per bug remediation due to raw stack-trace and framework bloat, costing \$2 to \$5 per issue while regularly generating broken diff headers and indentation syntax errors that cause CI/CD build failures.

[QUESTION B: What is your product and solution?]

Kronumos is an autonomous program repair (APR) engine engineered on a Dual-Brain cybernetic architecture. We pair neural reasoning models (Kronumos Kairos 7B/14B and Kronumos Aion 671B MoE) with an unyielding, zero-allocation native Rust Sub-Cortex compiler kernel (C-ABI 5µs). The Sub-Cortex deterministically excises 93.5% of issue context bloat (averaging ~2,500 tokens per task) and validates AST syntax via Tree-sitter before any modification touches production git branches, guaranteeing zero dirty diffs on Princeton SWE-bench Verified.

[QUESTION C: Why do you need Microsoft Azure compute?]

We require Azure compute credits to host and serve our flagship enterprise model, Kronumos Aion (671B MoE), on multi-GPU Azure Virtual Machine clusters (8x A100/H100 nodes via vLLM). Azure credits will also power our continuous integration benchmark testbed across the full 500-instance Princeton SWE-bench Verified dataset and support our upcoming automated GitHub Actions remediation bot for enterprise engineering teams.

3. FOUNDING TEAM PROFILE & ROLES

-
1. Muhammad Naufal Daffa – Chief Technology Officer (CTO) & Chief Scientist
Role: Architect of the Tokenectomy Rust Sub-Cortex, author of the Springer Nature preprint, and lead AI/ML systems engineer.
ORCID: <https://orcid.org/0009-0000-7909-4916>
 2. Muhammad Sufyan – Chief Executive Officer (CEO) & Head of Business Operations
Role: Overseeing commercial strategy, enterprise B2B pilot partnerships, grant applications, legal structure, and day-to-day operations.
 3. Adi Supriyadi – Head of Product & Lead Frontend Architect
Role: Leading developer experience, web portals, client monitoring dashboards, IDE extensions (VS Code / Cursor), and visual product design.

4. OPERATIONAL MILESTONES (FOR PROPOSAL REFERENCE)

-
- * Community Traction: 1,260+ organic model downloads on Hugging Face.
 - * Empirical Benchmark: Evaluated against all 500 Princeton SWE-bench Verified tasks.
 - * Sub-Cortex Efficiency: Verified 93.5% input token reduction at 5-microsecond latency.
 - * Academic Recognition: Peer-reviewed preprint published on Springer Nature Research Square and archived with CERN Zenodo.
-
-